


# Chapter 7

## Space and Time Tradeoffs



Copyright © 2007 Pearson Addison-Wesley. All rights reserved.



# Space-for-time tradeoffs

Two varieties of space-for-time algorithms:

- ▶ input enhancement — preprocess the input (or its part) to store some info to be used later in solving the problem
  - counting sorts
  - string searching algorithms
- ▶ prestructuring — preprocess the input to make accessing its elements easier
  - hashing
  - indexing schemes (e.g., B-trees)

# Review: String searching by brute force

*pattern*: a string of  $m$  characters to search for

*text*: a (long) string of  $n$  characters to search in

## Brute force algorithm

**Step 1** Align pattern at beginning of text

**Step 2** Moving from left to right, compare each character of pattern to the corresponding character in text until either all characters are found to match (successful search) or a mismatch is detected

**Step 3** While a mismatch is detected and the text is not yet exhausted, realign pattern one position to the right and repeat Step 2

# Brute-force string matching

```
for (i=0; T[i] != '\0'; i++)  
{  
  for (j=0; T[i+j] != '\0' && P[j] != '\0' && T[i+j]==P[j]; j++) ;  
  if (P[j] == '\0') found a match  
}
```

# String searching by preprocessing

Several string searching algorithms are based on the input enhancement idea of **preprocessing the pattern**

- ▶ **Knuth-Morris-Pratt (KMP)** algorithm preprocesses pattern **left to right** to get useful information for later searching
- ▶ **Boyer -Moore** algorithm preprocesses pattern **right to left** and store information into two tables
- ▶ **Horspool's** algorithm **simplifies** the Boyer-Moore algorithm by using just one table

# REF: Knuth-Morris-Pratt (KMP)

<http://www.ics.uci.edu/~eppstein/161/960227.html>

0 1 2 3 4 5 6 7 8 9 10 11

T: b a n a n a n o b a n o

i=0: X

i=1: X

i=2: n a n X

i=3: X

i=4: n a n o

i=5: X

i=6: n X

i=7: X

i=8: X

i=9: n X

i=10: X

# KMP

- ▶ String matching with skipped outer & inner iterations: KMP, version 1:

```
i=0;
o=0;
while (i<n)
{
  for (j=o; T[i+j] != '\0' && P[j] != '\0' && T[i+j]==P[j]; j++) ;
  if (P[j] == '\0') found a match;
  o = overlap(P[0..j-1],P[0..m]);
  i = i + max(1, j-o);
}
```

# KPM

## ► KMP, version 2:

```
j = 0;
for (i = 0; i < n; i++)
  for (;;) {    // loop until break
    if (T[i] == P[j]) { // matches?
      j++;      // yes, move on to next state
      if (j == m) { // maybe that was the last state
        found a match;
        j = overlap[j];
      }
      break;
    } else if (j == 0) break; // no match in state j=0, give up
    else j = overlap[j]; // try shorter partial match
  }
```



# KPM

## ► KMP overlap computation:

```
overlap[0] = -1;
for (int i = 0; pattern[i] != '\0'; i++) {
    overlap[i + 1] = overlap[i] + 1;
    while (overlap[i + 1] > 0 &&
           pattern[i] != pattern[overlap[i + 1] - 1])
        overlap[i + 1] = overlap[overlap[i + 1] - 1] + 1;
}
return overlap;
```

**Overlap:**  $O(m)$

**Scan:**  $O(n)$  time

**Total:**  $O(m+n)$ .

# Horspool's Algorithm

A simplified version of Boyer-Moore algorithm:

- preprocesses pattern to generate a shift table that determines how much to shift the pattern when a mismatch occurs
- always makes a shift based on the text's character  $c$  aligned with the last character in the pattern according to the shift table's entry for  $c$

# How far to shift?

Look at first (rightmost) character in text that was compared:

- ▶ The character is not in the pattern

.....**C**..... (C not in pattern)

BAOBAB

- ▶ The character is in the pattern (but not the rightmost)

.....**O**..... (O occurs once in pattern)

BA**O**BAB

.....**A**..... (A occurs twice in pattern)

BA**A**OBAB

- ▶ The rightmost characters do match

.....**B**.....

BAOBAB

# Shift table

- ▶ Shift sizes can be precomputed by the formula  
distance from  $c$ 's rightmost occurrence in pattern  
among its first  $m-1$  characters to its right end

$$t(c) =$$

pattern's length  $m$ , otherwise

**by scanning pattern before search begins and stored in a table called *shift table***

- Shift table is indexed by text and pattern alphabet  
Eg, for **BAOBAB**:

[illegible]

# Example of Horspool's alg. application

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	—
1	2	6	6	6	6	6	6	6	6	6	6	6	6	3	6	6	6	6	6	6	6	6	6	6	6	6

BARD LOVED BANANAS

BAOBAB (L=6)

BAOBAB (B=2)

BAOBAB (N=6)

BAOBAB (unsuccessful search)

PINGADO em

ALMOCEI PINGA COM LINGUADO PINGADO

▶ PINGADO outras

▶ 6543217 7

▶ 1234567890123456789012345678901234

▶ ALMOCEI PINGA COM LINGUADO PINGADO

▶ PINGADO (I=5)

▶ - - - - - PINGADO (^G=3)

▶ - - - PINGADO (^C=7)

▶ - - - - - PINGADO (^G=3)

▶ - - - PINGADO (^D=1)

▶ - PINGADO (^=ADO, U)

▶ PINGADO (^U=7)

▶ - - - - - PINGADO (^D=1)

▶ - PINGADO (=PINGADO)


# Boyer-Moore algorithm

Based on same two ideas:

- comparing pattern characters to text from right to left
- precomputing shift sizes in two tables
  - *bad-symbol table* indicates how much to shift based on text's character causing a **mismatch**
  - *good-suffix table* indicates how much to shift based on **matched** part (suffix) of the pattern

# Bad-symbol shift in Boyer-Moore algorithm

- ▶ If the rightmost character of the pattern **doesn't match**, BM algorithm acts as Horspool's
- ▶ If the rightmost character of the pattern **does match**, BM compares **preceding characters** right to left until either all pattern's characters match or a mismatch on text's character ***c*** is encountered after ***k* > 0** matches



bad-symbol shift  $d_1 = \max\{t_1(c) - k, 1\}$



# Good-suffix shift in Boyer-Moore algorithm

- ▶ Good-suffix shift  $d_2$  is applied after  $0 < k < m$  last characters were matched
- ▶  $d_2(k)$  = the distance between matched suffix of size  $k$  and its rightmost occurrence in the pattern that is not preceded by the same character as the suffix

Example: CABABA  $d_2(1) = 4$

- ▶ If there is no such occurrence, match the longest part of the  $k$ -character suffix with corresponding prefix; if there are no such suffix-prefix matches,  $d_2(k) = m$

Example: WOWWOW  $d_2(2) = 5$ ,  $d_2(3) = 3$ ,  $d_2(4) = 3$ ,  $d_2(5) = 3$

# Boyer-Moore Algorithm

After matching successfully  $0 < k < m$  characters, the algorithm shifts the pattern right by

$$d = \max \{d_1, d_2\}$$

where  $d_1 = \max\{t_1(c) - k, 1\}$  is bad-symbol shift

$d_2(k)$  is good-suffix shift

Example: Find pattern **AT\_THAT** in

**WHICH\_FINALLY\_HALTS. \_\_AT\_THAT**

# Boyer-Moore Algorithm (cont.)

Step 1 Fill in the bad-symbol shift table

Step 2 Fill in the good-suffix shift table

Step 3 Align the pattern against the beginning of the text

Step 4 Repeat until a matching substring is found or text ends:

Compare the corresponding characters right to left.

If no characters match, retrieve entry  $t_1(c)$  from the bad-symbol table for the text's character  $c$  causing the mismatch and shift the pattern to the right by  $t_1(c)$ .

If  $0 < k < m$  characters are matched, retrieve entry  $t_1(c)$  from the bad-symbol table for the text's character  $c$  causing the mismatch and entry  $d_2(k)$  from the good-suffix table and shift the pattern to the right by

$$d = \max \{d_1, d_2\}$$

where  $d_1 = \max\{t_1(c) - k, 1\}$ .

# Example of Boyer-Moore alg. application

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W	X	Y	Z	—
1	2	6	6	6	6	6	6	6	6	6	6	6	6	3	6	6	6	6	6	6	6	6	6	6	6	6

B E S S \_ **K** N E W \_ A B O U T \_ B A O B A B S

B A O B A **B**

$d_1 = t_1(\mathbf{K}) = 6$       B A O **B** A B

$d_1 = t_1(\_) - 2 = 4$

$d_2(2) = 5$

*$\max\{d_1, d_2\}$*

$k$	pattern	$d_2$
1	BAO <b>B</b> AB	2
2	<b>B</b> AOBAB	5
3	<b>B</b> AOBAB	5
4	<b>B</b> AOBAB	5
5	<b>B</b> AOBAB	5

B A O B A B

$d_1 = t_1(\_) - 1 = 5$

$d_2(1) = 2$

B A O B A B (success)

# Hashing

- ▶ A very efficient method for implementing a *dictionary*, i.e., a set with the operations:
  - find
  - insert
  - delete
- ▶ Based on **representation-change** and **space-for-time** tradeoff ideas
- ▶ Important applications:
  - symbol tables
  - databases (*extendible hashing*)

# Hash tables and hash functions

The idea of *hashing* is to map keys of a given file of size  $n$  into a table of size  $m$ , called the *hash table*, by using a predefined function, called the *hash function*,

$h: K \rightarrow \text{location (cell) in the hash table}$

Example: student records, key = **SSN**. Hash function:

$h(K) = K \bmod m$  where  $m$  is some integer (typically, **prime**)

If  $m = 1000$ , where is record with SSN= 314159265 stored?

Generally, a hash function should:

- be easy to compute
- distribute keys about evenly throughout the hash table

# Collisions

If  $h(K_1) = h(K_2)$ , there is a *collision*

- ▶ Good hash functions result in fewer collisions but some collisions should be expected (*birthday paradox*)
- ▶ Two principal hashing schemes handle collisions differently:
  - *Open hashing*
    - each cell is a header of **linked list** of all keys hashed to it
  - *Closed hashing*
    - one key per cell
    - in case of collision, finds another cell by
      - *linear probing*: use **next** free bucket
      - *double hashing*: use **second hash** function to compute increment


# Open hashing (Separate chaining)

Keys are stored in linked lists outside a hash table whose elements serve as the lists' headers.

Example: A, FOOL, AND, HIS, MONEY, ARE, SOON, PARTED

$h(K)$  = sum of  $K$ 's letters' positions in the alphabet MOD 13

Key	A	FOOL	AND	HIS	MONEY	ARE	SOON	PARTED
$h(K)$	1	9	6	10	7	11	11	12



Example: Search for KID



# Open hashing (cont.)

- ▶ If hash function distributes keys uniformly, average length of linked list will be  $\alpha = n/m$ . This ratio is called *load factor*.
  - file of size  $n$  into a table of size  $m$

- ▶ Average number of probes in successful,  $S$ , and unsuccessful searches,  $U$ :

$$S \approx 1 + \alpha/2, \quad U = \alpha$$

- ▶ Load  $\alpha$  is typically kept small (ideally, about 1)
- ▶ Open hashing still works if  $n > m$

# Closed hashing (Open addressing)

Keys are stored inside a hash table.

Key	A	FOOL	AND	HIS	MONEY	ARE	SOON	PARTED
$h(K)$	1	9	6	10	7	11	11	12

	0	1	2	3	4	5	6	7	8	9	10	11	12
		A											
		A								FOOL			
		A					AND			FOOL			
		A					AND			FOOL	HIS		
		A					AND	MONEY		FOOL	HIS		
		A					AND	MONEY		FOOL	HIS	ARE	
		A					AND	MONEY		FOOL	HIS	ARE	SOON
→ PARTED		A					AND	MONEY		FOOL	HIS	ARE	SOON

# Closed hashing (cont.)

- ▶ Does not work if  $n > m$ 
  - file of size  $n$  into a table of size  $m$
- ▶ Avoids pointers
- ▶ Deletions are *not* straightforward
- ▶ Number of probes to find/insert/delete a key depends on load factor  $\alpha = n/m$  (hash table density) and collision resolution strategy. For linear probing:  
$$S = (1/2) (1 + 1/(1 - \alpha)) \text{ and } U = (1/2) (1 + 1/(1 - \alpha)^2)$$
- ▶ As the table gets filled ( $\alpha$  approaches 1), number of probes in linear probing increases dramatically:

$\alpha$	$\frac{1}{2} \left( 1 + \frac{1}{1-\alpha} \right)$	$\frac{1}{2} \left( 1 + \frac{1}{(1-\alpha)^2} \right)$
50%	1.5	2.5
75%	2.5	8.5
90%	5.5	50.5

# indexing schemes (e.g., B-trees)

- ▶ **Idea**
- ▶ **Example**
- ▶ **Height**
- ▶ **B\* : redistribuição**
- ▶ **B+ : chaves nas folhas**